

# A Practical Tutorial for Jupyter Notebooks

Extracted from Logan Ward's notes

Postdoctoral Scholar

University of Chicago

17 October 2018

# Materials Informatics Workflow



2

Collect

Process

Represent

Learn



$$\Delta H_f = -1.0$$

$$\Delta H_f = -0.5$$

$\vec{X}$

$\vec{y}$

$Z_A$	$Z_B$	$\Delta H_f$
3	4	-1.0
3	5	-0.5

$$\Delta H_f = f(Z_A, Z_B)$$

# Data Representation

3

Collect


Process

Represent

Learn


***Encode domain knowledge into inputs***

$$\text{Property} = f(\text{Attributes})$$



Representation of material

$$\text{Ex: } \text{Attributes} = \mathbf{g}(x_H, x_{He}, \dots)$$



**What does a representation need?**

*Completeness:* Differentiate materials

*Efficiency:* Quick to compute

*Accuracy:* Capture important effects

**End product:** Machine-learning compatible inputs

# Fit a model to the data

4

Collect

Process

Represent

Learn

**There are many algorithms, how do you decide?**

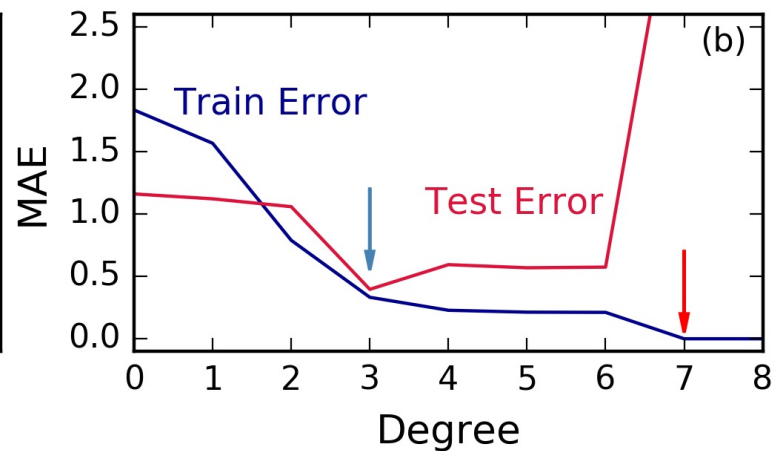

1. Identify algorithms with desired properties.

Ex: differentiability

2. Use cross-validation to optimize parameters

3. Pick the one with the lowest “loss”

- Numpy
- Scipy
- Scikit-learn
- Theano
- TensorFlow
- Keras
- PyTorch
- Pandas
- Matplotlib



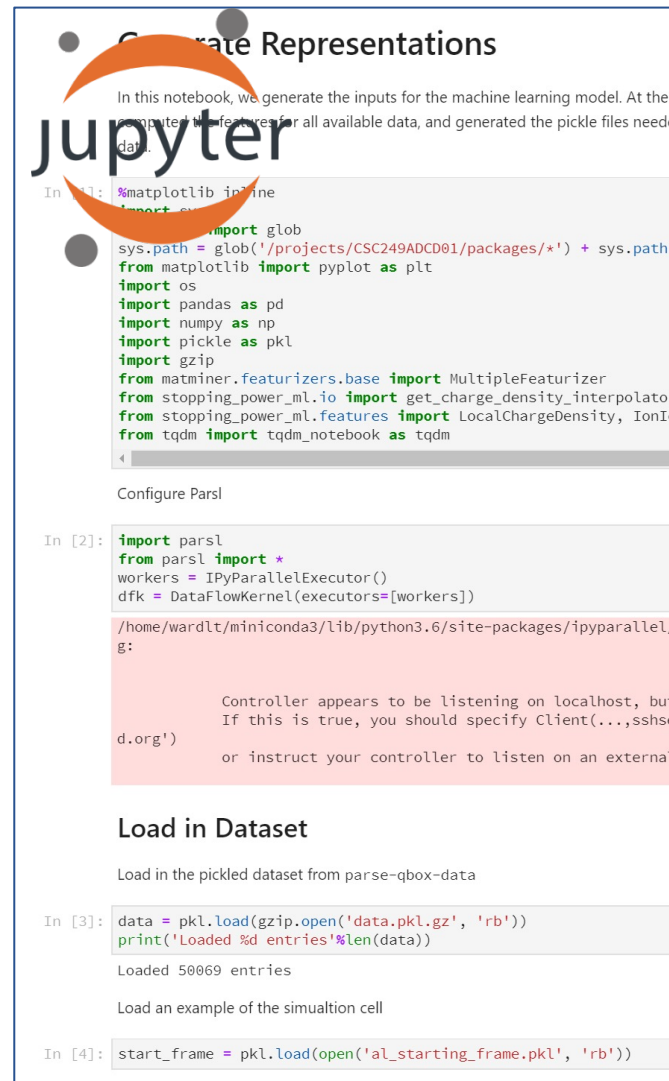
5

# How to Organize? [Jupyter]

# What is Jupyter? Why care?

6

- Jupyter is environment designed for reproducible computational science
- Stores code with outputs and documentation in a single notebook
  - Modeled after Mathematica
  - <https://www.theatlantic.com/science/archive/2018/04/the-scientific-paper-is-obsolete/556676/>
- Why do I care? Jupyter lets me...
  - organize my code
  - easily run on a remote system
  - keep track of results and rationale
  - communicate my research better



The screenshot shows a Jupyter Notebook with the following content:

- Generate Representations**  
In this notebook, we generate the inputs for the machine learning model. At the computer the features for all available data, and generated the pickle files need data.
- In [1]:**  

```
%matplotlib inline
import sys
import glob
sys.path = glob('/projects/CSC249ADC001/packages/*') + sys.path
from matplotlib import pyplot as plt
import os
import pandas as pd
import numpy as np
import pickle as pkl
import gzip
from matminer.featurizers.base import MultipleFeaturizer
from stopping_power_ml.io import get_charge_density_interpolator
from stopping_power_ml.features import LocalChargeDensity, IonI
from tqdm import tqdm_notebook as tqdm
```
- Configure Parsl**  
**In [2]:**  

```
import parsl
from parsl import *
workers = IPyParallelExecutor()
dfk = DataFlowKernel(executors=[workers])

/home/wardlt/miniconda3/lib/python3.6/site-packages/ipyparallel/
g:
```

Controller appears to be listening on localhost, but  
If this is true, you should specify Client(...,sshse  
d.org')

or instruct your controller to listen on an external
- Load in Dataset**  
Load in the pickled dataset from parse-qbox-data  
**In [3]:**  

```
data = pkl.load(gzip.open('data.pkl.gz', 'rb'))
print('Loaded %d entries'%len(data))
```

Loaded 50069 entries

Load an example of the simulation cell

**In [4]:**  

```
start_frame = pkl.load(open('al_starting_frame.pkl', 'rb'))
```

# How I write a notebook

7

*One notebook per “experiment” or “idea”*

## Outline:

1. Title and short abstract
  1. “What am I doing here and why”
2. Load in libraries
  1. Put them up front, so it crashes early
3. Load in data from disk
  1. Ex: training data, results from other notebook
4. Each step in their own block
  1. Introduction
  2. Code and explanation
  3. Figure/visualization
  4. Explanation of finding

# Example Step

8

## Simple Test: Channel

### Introduction

Here, we have a particle traveling forward or backwards along the channel of FCC AI. The path forward and backwards are identical, so we should get the same stopping power

```
In [5]: @App('python', dfk)
def compute_stopping_power(starting_point, direction, traj_computer=traj_computer):
    return traj_computer.compute_stopping_power(starting_point, direction, 1)
```

Set up calculations for the stopping power in the channel

```
In [6]: forward = compute_stopping_power([0,0.75,0.75], [1,0,0])
```

```
In [7]: backward = compute_stopping_power([0,0.75,0.75], [-1,0,0])
```

Wait for them to finish

### Explanation

```
In [8]: %%time
forward = forward.result(); backward = backward.result()
```

CPU times: user 12 ms, sys: 12 ms, total: 24 ms  
Wall time: 1.8 s

```
In [9]: print('Forward stopping power: ', forward[0])
print('Backward stopping power: ', backward[0])
print('Difference: ', backward[0]-forward[0])
```

Forward stopping power: 0.2343000208236084  
Backward stopping power: 0.23431760358559353  
Difference: 1.758276198512987e-05

### Visualization and conclusion

*Finding:* Consistent with my initial expectations, they are indeed the same (within numerical tolerances).

Someone should be able to understand this without knowing Python!



# The Full Narrative

9

## Assessing Forward/Backward Asymmetry in Stopping Powers

Our stopping power model predicts different stopping powers for particles traveling forward than those traveling backwards on some trajectories. The purpose of this notebook is to showcase why this is valid.

```
In [1]: %matplotlib inline
from matplotlib import pyplot as plt
from stopping_power_ml.features import LocalChargeDensity
import pandas as pd
import numpy as np
import pickle as pkl
import os
```

Configure parsl

```
In [2]: import parsl
from parsl import *
#from parsl_config import config
workers = IPyParallelExecutor()
dfk = DataFlowKernel(executors=[workers])
print("Parsl version : ", parsl.__version__)
```

Parsl version : 0.5.0

/home/wardlt/miniconda3/lib/python3.6/site-packages/ipyparallel/client/client.py:458: RuntimeWarning:

Controller appears to be listening on localhost, but not on this machine.  
If this is true, you should specify Client(...,sshserver='you@js-168-224.jetstream-cloud.org')  
or instruct your controller to listen on an external IP.

## Load in the Tools

We'll need the trajectory computer, and the charge density so that we can make illustrative plots.

```
In [3]: traj_computer = pkl.load(open('traj_computer.pkl', 'rb'))
```

```
In [4]: charge_density = pkl.load(open(os.path.join('.', 'density_interp.pkl'), 'rb'))
```

[Link to GitHub Page](#)

# Pitfalls When Making Notebooks

10

1. Not including documentation

**Advice:** Write what you're going to do first

2. Notebook not capturing entire process

**Problem:** Jupyter lets you execute cells out of order

**Advice:** Periodically “Restart Kernel and Run All Cells”

3. Duplicate code between notebooks

**Advice:** Make a separate module for common code

4. Library conflicts

**Advice:** Run mature projects in container, separate machines

**Advice:** Make a “requirements.txt” file

5. The “one cell notebooks”

**Advice:** <10 lines of code per cell

See “I Don’t Like Jupyter Notebooks”: <https://t.co/30peBFwTbv>

# Organizing Multiple Notebooks

11

🏠 > single-velocity

Name	Last Modified
📁 convergence-detection	a month ago
📁 data	a month ago
📁 feature-analysis	a month ago
📁 figures	a month ago
📁 manifold	17 hours ago
📁 neural-network	15 hours ago
📁 runinfo	15 hours ago
📄 0_collect-subset.ipynb	16 hours ago
📄 1_build-machine-learni...	16 hours ago
📄 2_evaluate-direction-d...	16 hours ago
📄 3_forward-backward-as... in a few seconds	
📄 best_model.pkl	16 hours ago
📄 best_weight.pkl	a month ago
📄 direction_stopping.pkl	16 hours ago
📄 run-notebooks.sh	a month ago
📄 stopping_power_result...	16 hours ago
📄 traj_computer.pkl	16 hours ago

Re-usable datasets in separate folder

Figures in separate folder  
(with meaningful names!)  
(making pub-ready figures!)

Dependent tasks in separate folder

Notebooks numbered by  
execution order

Results shared between  
steps using pickled files

A bash script to re-run all notebooks

# Running Notebooks from Terminal

12

```
#!/bin/bash
```

```
# This script just runs through all of the current notebooks in the proper order
```

```
runJupyter() {  
    jupyter nbconvert --execute --inplace --ExecutePreprocessor.timeout=-1 $1  
}
```

```
# If we get too many files, we should probably create a file containing  
# names of notebooks and the order in which they should be executed  
for n in `seq 0 1`; do  
    runJupyter ${n}_*nb  
done
```

Rerun your project often, put in “assert” statements for key results!

# Keeping Track of Project

13

Commits on May 21, 2018

Switched to using spherical coordinates for traj optimizer

WardLT committed 3 days ago

**Use GitHub to track changes**

Commits on May 20, 2018

Use global optimizer to find best trajectory ...

WardLT committed 4 days ago



d027633



Overview

Yours

Active

Stale

All branches

Search branches...

Default branch

master Updated 3 days ago by WardLT

**Make branches for dead-ends**

Your branches

Branch: master ▾

stopping-power-ml / 1\_compute-representation.ipynb

Find file

Copy path

WardLT Switched back to using time offsets for the change density

6d7d3f2 7 days ago

1 contributor

459 lines (458 sloc) 86.5 KB

**GitHub renders your notebooks:  
Great for Sharing with Collaborators!**

## Generate Representations

In this notebook, we generate the inputs for the machine learning model. At the end of this notebook, we will have computed the features for all available data, and generated the pickle files needed to run these models for any available data.

# Take-Away Points

14

1. Write your notebooks like papers  
*Explain what you are doing, why, and what you found*  
*Learn Markdown to include links, equations, pictures*
2. Break up complex projects into multiple steps  
*Each notebook should tell a single story*
3. Notebooks should be easy to re-run  
*Use “Restart and Run,” bash scripts*
4. Track your changes with Git  
*Explain what changes you made, and why*

15

# Publication: Not Just Papers

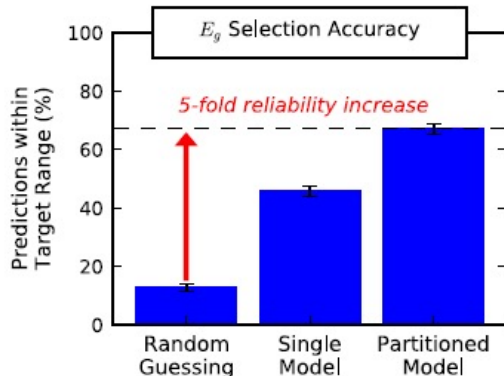
# Is my 2016 paper reproducible?

16

## ARTICLE OPEN

# A general-purpose machine learning framework for predicting properties of inorganic materials

Logan Ward<sup>1</sup>, Ankit Agrawal<sup>2</sup>, Alok Choudhary<sup>2</sup> and Christopher Wolverton<sup>1</sup>



**Figure 1.** Performance of three different strategies to locate compounds with a band gap energy within a desired range: randomly selecting nonmetal-containing compounds, and two based methods presented in this work. The Partitioned Model strategy used a single model trained on the band gap energies of 22,667 compounds.

**Table 2.** Compositions and predicted band gap energies of materials predicted using machine learning to be candidates for solar cell applications

Composition	$E_g$ (eV)
ScHg <sub>4</sub> Cl <sub>7</sub>	1.26
V <sub>2</sub> Hg <sub>3</sub> Cl <sub>7</sub>	1.16
Mn <sub>6</sub> CCl <sub>8</sub>	1.28
Hf <sub>4</sub> S <sub>11</sub> Cl <sub>2</sub>	1.11
VCu <sub>5</sub> Cl <sub>9</sub>	1.19

Abbreviations: DFT, density functional theory; OQMD, open quantum materials database.

Compositions represent the nominal compositions of novel ternary materials developed in ref. 15. Band gap energies were predicted using a machine learning model trained on DFT calculations and the methods described in this work.

Method Comparison

Predicted Materials



# I released the code, but not *all* of it

17

## From Original Source Code


### Supplementary information

PDF files

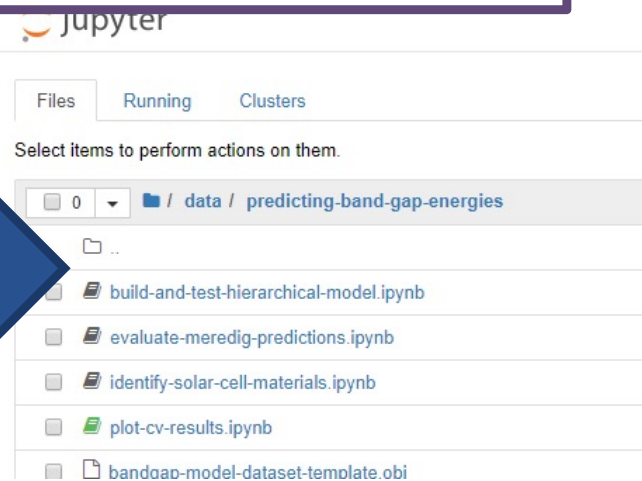
1. Supplementary Information

Zip files

1. Supplementary Information




## ... to New Notebooks




# ... and there are differences

18

## Reported Results in 2016(?)



## Replication in 2017



**Without Exact Scripts From <2016, Cannot Verify Results**

predicted using machine learning to be candidates for solar cell applications

Out[11]:

Composition

SCHg <sub>4</sub> Cl <sub>7</sub>	1.1
V <sub>2</sub> Hg <sub>3</sub> Cl <sub>7</sub>	1.28
Mn <sub>6</sub> CCl <sub>8</sub>	1.11
Hf <sub>4</sub> S <sub>11</sub> Cl <sub>2</sub>	
VCu <sub>5</sub> Cl <sub>9</sub>	

Abbreviations: DF  
materials database



	Entry	bandgap_predicted
1037	CoB2F9	1.380256
3414	YbAs7Cl6	1.156973
3884	Tl3OsO3.5	1.078918
1920	Cs8CoSe5	1.074358
		1.119002

**Publication Should Go Beyond the Paper!**

# How Do We Avoid this Problem?

19

## What do I need to publish?

- *Datasets*: In a well-described format
- *Scripts*: Not just the core methods
- *Outputs*: Exact version from the paper
- *Models*: In a user-friendly way

Save Anaconda/Python  
versions and Environments/  
library version!

